

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix

df=pd.read_csv(r"C:\Users\ASUS\Documents\pythonStack\DS_PR\
Social_Network_Ads.csv")

df.head()

```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19.0	19000.0	0
1	15810944	Male	35.0	20000.0	0
2	15668575	Female	26.0	43000.0	0
3	15603246	Female	27.0	57000.0	0
4	15804002	Male	19.0	76000.0	0

```

xtrain=df[['Age', 'EstimatedSalary']]
ytrain=df['Purchased']

xtrain,xtest,ytrain,ytest=train_test_split(xtrain,ytrain,test_size=0.2
)

xtrain

```

	Age	EstimatedSalary
30	31.0	74000.0
300	58.0	38000.0
88	26.0	81000.0
282	37.0	70000.0
181	31.0	71000.0
..
64	59.0	83000.0
230	35.0	147000.0
49	31.0	89000.0
227	56.0	133000.0
183	33.0	43000.0

```

[304 rows x 2 columns]

model=LogisticRegression()

model.fit(xtrain,ytrain)

LogisticRegression()

yprediction=model.predict(xtest)

ytest

```

244	0
377	0

```
166    0
110    0
94     0
```

```
..
233    1
170    0
123    0
386    1
138    0
```

```
Name: Purchased, Length: 76, dtype: int64
```

```
CM=confusion_matrix(ytest,yprediction)
```

```
CM
```

```
array([[48,  4],
       [ 1, 23]], dtype=int64)
```

```
# Confusion Matrix:
```

```
#                Predicted: 0    Predicted: 1
```

```
# Actual: 0      48 (TN)         4 (FP)
```

```
# Actual: 1      1 (FN)         23 (TP)
```

```
#
```

```
# TN = True Negative → predicted 0, actually 0
```

```
# FP = False Positive → predicted 1, actually 0
```

```
# FN = False Negative → predicted 0, actually 1
```

```
# TP = True Positive → predicted 1, actually 1
```

```
tn, fp, fn, tp = CM.ravel()
```

```
#unpacks the values from the confusion matrix into individual variables
```

```
#.ravel() flattens a 2D array into a 1D list[0,0,0,0]
```

```
tn,fp,fn,tp
```

```
(48, 4, 1, 23)
```

```
#Purpose: Measures how often the model predicted correctly overall.
```

```
accuracy = (tp + tn) / (tp + tn + fp + fn)
```

```
#Because 1 means 100% of predictions, subtracting accuracy gives incorrect percentage.
```

```
#Shows how often the model made wrong predictions.
```

```
error_rate = 1 - accuracy
```

```
accuracy.round(),error_rate.round()
```

```
(1.0, 0.0)
```

```
Psion= tp / (tp + fp)
```

```
#Purpose: Of all predicted 1s, how many are actually 1?
```

*#Numerator: tp → Correctly predicted positives.
#Denominator: tp + fp
#fp: False Positives (predicted 1, but actually 0)*

Psion

0.8518518518518519

*recall = tp / (tp + fn)
Recall tells how many actual positives the model correctly predicted.
Formula: Recall = TP / (TP + FN)
Used when missing positive cases is risky (e.g., disease detection, fraud).
High recall means fewer false negatives (model didn't miss many real positives).
Example: If 10 people actually bought and model found 8 → Recall = 8/10 = 0.8
#80 percentr real buyer found by model*

recall

0.9583333333333334